Smile Dynamics II

Lorenzo Bergomi
Société Générale
lorenzo.bergomi@sgcib.com

March 2005

Abstract

In a previous article we highlighted how traditional stochastic volatility and Jump/Lévy models impose structural constraints on how the short forward skew, the spot/vol correlation, and the term-structure of the vol-of-vol are related. Here we propose a model that enables them to be controlled separately and also prices options on realized variance consistently. We present pricing examples for a reverse cliquet, a Napoleon, an accumulator and an option on variance.

1. Introduction

A common feature of the recent breed of exotic options such as Napoleons and reverse cliquets is that their price depends on assumptions made for the joint dynamics of the underlying and its implied volatilities. These fall into three categories:

- The dynamics of implied volatilities, and more specifically the term-structure of the volatility of volatility.
- The forward skew.
- The spot/vol correlation.

In a previous paper (Bergomi 2004), we analyzed popular stochastic volatility and Jump/Lévy models and pointed out that, although these models produce prices which include an estimation of the three effects listed above, they impose structural constraints on how these features of the joint dynamics of the spot and implied volatilities are related. Another of their drawbacks is that they are based on a specification of the spot process and fail to take into account the fact that Variance Swaps (VS) should be considered as hedge instruments too, and be endowed with their own dynamics.

This article is a natural continuation of our first one: in this issue we propose a model which aims at pricing both standard exotic options and general options on variance in a consistent manner, and lets us independently set requirements on:

- The dynamics of VS volatilities
- The level of short-term forward skew
- The correlation between the underlying and short and long VS volatilities

The article is organized as follows: we first set up a general framework for the dynamics of Forward Variance Swap Variances - which we call simply "variances". Then we specify a dynamics for the underlying which is consistent with that of variances. In the next section we specify a particular choice for the dynamics of FVs and the underlying. The following section focuses on practical features of the model such as the term-structure of the volatility of volatility an the term-structure of the skew. Then a section focuses on pricing examples: we consider a reverse cliquet, a Napoleon, an accumulator and a call on variance. The concluding section summarizes our work.

2. Modelling variances

A VS pays at maturity $V_{tT}^h - V_t^T$ where V_{tT}^h is the annualized variance of the spot, realized over the interval $[t, T]$ and V_t^T is the implied VS variance, observed at time t for maturity T. Because VSs are statically replicable by vanilla options V_t^T only depends on the implied vols seen at time t for maturity T[1]. Because of the definition of V_t^T the VS contract has zero value at inception.

Now consider the variance $V_t^{T_1,T_2}$ defined as:

$$V_t^{T_1,T_2} = \frac{(T_2 - t)V_t^{T_2} - (T_1 - t)V_t^{T_1}}{T_2 - T_1}$$

where $T_1, T_2 > t$.

To write a pricing equation for an option on $V_t^{T_1,T_2}$ we first need to know the cost of entering a trade whose payoff at time $t+dt$ is linear in $V_{t+dt}^{T_1,T_2} - V_t^{T_1,T_2}$. Let us buy $\frac{T_2-t}{T_2-T_1}e^{r(T_2-t)}$VS of maturity T_2 and sell $\frac{T_1-t}{T_2-T_1}e^{r(T_1-t)}$VS of maturity T_1. This is done at no cost; our P&L at time $t' = t + dt$ is:

$$\begin{aligned} P\&L &= \frac{T_2 - t}{T_2 - T_1}\left(\frac{V_{tt'}^h\,(t'-t) + V_{t'}^{T_2}\,(T_2 - t')}{T_2 - t} - V_t^{T_2}\right)e^{r(T_2-t)}e^{-r\left(T_2-t'\right)} \\ &\quad - \frac{T_1 - t}{T_2 - T_1}\left(\frac{V_{tt'}^h\,(t'-t) + V_{t'}^{T_1}\,(T_1 - t)}{T_1 - t} - V_t^{T_1}\right)e^{r(T_1-t)}e^{-r\left(T_1-t'\right)} \\ &= \left(V_{t'}^{T_1,T_2} - V_t^{T_1,T_2}\right)e^{-r\left(t'-t\right)} = \left(V_{t+dt}^{T_1,T_2} - V_t^{T_1,T_2}\right)(1 - rdt) \end{aligned}$$

This position generates a P&L linear in $V_{t+dt}^{T_1,T_2} - V_t^{T_1,T_2}$ at lowest order in dt, at zero initial cost: thus the pricing drift of any forward FV $V_t^{T_1,T_2}$ is zero[2].

We now specify a dynamics for the VS curve. Let us introduce $\xi_t^T = V_t^{T,T}$, the value of the variance for date T, observed at time t.

2.1. A one-factor model

We are free to specify any dynamics on the $\xi^T(t)$ that complies with the requirement that $\xi^T(t)$ be driftless. However, for practical pricing purposes, we would like to drive the dynamics of all of the $\xi^T(t)$ with a small number of factors. In this paragraph we show how this can be done by carefully choosing the volatility function of $\xi^T(t)$.

Let us assume $\xi^T(t)$ is lognormally distributed and that its volatility is a function of $T-t$ so that the model is translationally invariant through time:

$$d\xi^T = \omega(T-t)\,\xi^T\,dU_t$$

where U_t is a Brownian motion. Let us choose the form $\omega(\tau) = \omega\, e^{-k_1\tau}$.

$\xi^T(t)$ can be written as:

$$\xi^T(t) = \xi^T(0)\,e^{\left(\omega e^{-k_1(T-t)}X_t - \frac{\omega^2}{2}e^{-2k_1(T-t)}E[X_t^2]\right)} \tag{2.1}$$

where X_t is an Ornstein-Ühlenbeck process

$$X_t = \int_0^t e^{-k_1(t-u)}dU_u$$

whose dynamics reads:

$$\begin{aligned} dX_t &= -k_1 X_t dt + dU_t \\ X_0 &= 0 \end{aligned}$$

$\xi^T(t)$ is driftless by construction. Knowing X_t, we can generate $X_{t+\delta}$ through:

$$X_{t+\delta} = e^{-k_1\delta}X_t + x_\delta$$

where x_δ is a centered Gaussian random variable such that $E[x_\delta^2] = \frac{1-e^{-2k_1\delta}}{2k_1}$.

[1] As well as on how dividends are modelled and assumptions on interest rate volatility.

[2] The driftless nature of forward VS variances had been noticed before - see Dupire (1996). In diffusive models it is dictated by the definition of forward variance as a martingale.

Starting from known values for X_t and $E\left[X_t^2\right]$ at time t we can generate the FV curve $\xi^T(t+\delta)$ at time $t+\delta$ by using the following relationship:

$$\begin{aligned} X_{t+\delta} &= e^{-k_1\delta}X_t + x_\delta \\ E\left[X_{t+\delta}^2\right] &= e^{-2k_1\delta}E\left[X_t^2\right] + \frac{1-e^{-2k_1\delta}}{2k_1} \end{aligned}$$

and expression (2.1).

Thus by choosing an exponentially decaying form for $\omega(\tau)$ the model becomes Markovian: all $\xi^T(t)$ are functions of just one Gaussian factor X_t.

2.2. A two-factor model

To achieve greater flexibility in the range of term-structures of volatilities of variances that can be generated, we prefer to work with two factors. We then write:

$$d\xi^T = \omega\xi^T\left(e^{-k_1(T-t)}dU_t + \theta e^{-k_2(T-t)}dW_t\right)$$

where W_t is a Brownian motion. Its correlation with U_t is ρ. We can run through the same derivation as above. $\xi^T(t)$ now reads:

$$\xi^T(t) = \xi^T(0)\exp\left(\begin{array}{c} \omega\left[e^{-k_1(T-t)}X_t + \theta e^{-k_2(T-t)}Y_t\right] \\ -\frac{\omega^2}{2}\left[e^{-2k_1(T-t)}E[X_t^2] + \theta^2 e^{-2k_2(T-t)}E[Y_t^2] + 2\theta e^{-(k_1+k_2)(T-t)}E[X_tY_t]\right] \end{array}\right) \tag{2.2}$$

As in the 1-factor case, if $X_t,\ Y_t,\ E\left[X_t^2\right],\ E\left[Y_t^2\right],\ E[X_tY_t]$ are known at time t, they can be generated at time $t+\delta$ through the following relationships:

$$\begin{aligned} X_{t+\delta} &= e^{-k_1\delta}X_t + x_\delta \\ Y_{t+\delta} &= e^{-k_2\delta}Y_t + y_\delta \end{aligned}$$

and

$$\begin{aligned} E\left[X_{t+\delta}^2\right] &= e^{-2k_1\delta}E\left[X_t^2\right] + \frac{1-e^{-2k_1\delta}}{2k_1} \\ E\left[Y_{t+\delta}^2\right] &= e^{-2k_2\delta}E\left[Y_t^2\right] + \frac{1-e^{-2k_2\delta}}{2k_2} \\ E\left[X_{t+\delta}Y_{t+\delta}\right] &= e^{-(k_1+k_2)\delta}E\left[X_tY_t\right] + \rho\frac{1-e^{-(k_1+k_2)\delta}}{k_1+k_2} \end{aligned}$$

where, in the right-hand side terms, the second component is, respectively, the variance of x_δ, the variance of y_δ and the covariance of x_δ and y_δ. Starting from time $t=0$ we can easily generate a FV curve at any future time t by simulating 2 Gaussian factors. We choose $k_1 > k_2$ and call X_t the short factor, Y_t the long factor.

2.3. A discrete structure

Instead of modelling the set of all instantaneous forward variances, it may be useful to set up a Tenor structure and model the dynamics of forward variances for discrete time intervals, in a way which is analogous to Libor market models.

In Fixed Income this is motivated by the fact that forward Libor rates are the actual underliers over which options are written. In our case, it is motivated by the fact that we want to control the skew for a given time scale.

Let us define a set of equally spaced dates $T_i = t_0 + i\Delta$, starting from t_0, today's date. We will model the dynamics of FVs defined over intervals of width Δ: define $\xi^i(t) = V_t^{t_0+i\Delta,\ t_0+(i+1)\Delta}$, for $t \le t_0 + i\Delta$. $\xi^i(t)$ is the value at time t of the FV for the interval $[t_0+i\Delta,\ t_0+(i+1)\Delta]$.

$\xi^i(t)$ is a random process until $t = t_0 + i\Delta$. When t reaches $t_0 + i\Delta$, the Variance Swap variance for time interval $[t, t+\Delta]$ is known and is equal to $\xi^i(t = t_0 + i\Delta)$.

We model the ξ^i in the same way as their continuous counterparts:

$$\xi^i(t) = \xi^i(0)\exp\left(\begin{array}{c} \omega\left[e^{-k_1(T_i-t)}X_t + \theta e^{-k_2(T_i-t)}Y_t\right] \\ -\frac{\omega^2}{2}\left[e^{-2k_1(T_i-t)}E[X_t^2] + \theta^2 e^{-2k_2(T_i-t)}E[Y_t^2] + 2\theta e^{-(k_1+k_2)(T_i-t)}E[X_tY_t]\right] \end{array}\right) \tag{2.3}$$

where we use the same recursions as above for $X_t,\ Y_t,\ E\left[X_t^2\right],\ E\left[Y_t^2\right],\ E[X_tY_t]$.

While this setup for the dynamics of the ξ^i is reminiscent of the Libor market models used in Fixed Income, there are as yet no market quotes for prices of caps/floors and swaptions on Forward Variances, on which to calibrate volatilities and correlations for the ξ^i.

2.4. An N-factor model

We may generally write

$$\xi^i(t) = \xi^i(0)\, e^{\omega_i Z_t^i - \frac{\omega_i^2 t}{2}}$$

where ω_i and $\rho(Z_i, Z_j)$ are chosen at will. Further in this article we will compare pricing results obtained in the two-factor model with those obtained in an N–factor model for which $\omega_i = \omega$, a constant, and the correlation structure of the Z^i is:

$$\rho(Z_i, Z_j) = \theta\rho_0 + (1-\theta)\,\beta^{|j-i|}. \tag{2.4}$$

where $\theta,\ \rho_0,\ \beta \in [0,1]$.

It should be noted that, when pricing an option of maturity T, in contrast to the 2-factor model, the number of factors driving the dynamics of variances in the N-factor model is proportional to T, thus the pricing time will grow like T^2.

3. Specifying a joint dynamics for the spot

3.1. A continuous setting

We could use the dynamics of instantaneous forward variances specified in eq. (2.2) and write the following lognormal dynamics on the underlying:

$$dS = (r-q)\,Sdt + \sqrt{\xi^t(t)}SdZ_t$$

with correlations ρ_{SX} and ρ_{SY} between Z and, respectively U and W. This yields a stochastic volatility model whose differences with standard models are:

- it has two factors
- it is calibrated by construction to the term-structure of VS volatilities

In such a model the level of forward skew is determined by $\rho_{SX},\ \rho_{SY}, \rho,\ \omega,\ k_1,\ k_2,\ \theta$ with no way of controlling it separately, just like in standard stochastic volatility models.

3.2. A discrete setting

Here we achieve our objective of controlling the forward skew – or, in other words, the skewness of the spot process for time scale Δ – by using the discrete Tenor structure defined above and the dynamics of forward variances given by expression (2.3) .

At time $t = T_i$, the VS volatility $\widehat{\sigma}_{VS}$ for maturity $T_i+\Delta$ is known: it is given by $\widehat{\sigma}_{VS} = \sqrt{\xi^i(T_i)}$. To be able to specify the spot process over the interval $[T_i,\ T_i+\Delta]$ we make a few more assumptions:

- the spot process over the time interval $[T_i,\ T_i+\Delta]$ is homogeneous: the distribution of $\frac{S_{T_i+\Delta}}{S_{T_i}}$ does not depend on S_{T_i}. The reason for this requirement is that we want to decouple the short forward skew and the spot/vol correlation. Imposing this condition makes the skew of maturity Δ independent on the spot level. Thus prices of cliquets of period Δ will not depend on the level of spot-vol correlation.
- we impose that the at-the-money-forward (ATMF) skew $\left.\frac{d\widehat{\sigma}_K}{d\ln K}\right|_F$ for maturity $T_i+\Delta$ be a deterministic function of $\widehat{\sigma}_{VS}$ or $\widehat{\sigma}_{ATMF}$. In this article we impose that it is constant or proportional to $\widehat{\sigma}_{ATMF}$. Other specifications for the dependence of the ATMF skew on $\widehat{\sigma}_{VS}$ or $\widehat{\sigma}_{ATMF}$ are easily accomodated in our framework.

There are many processes available for fulfilling our objective – note that we also need to correlate the spot process with that of forward variances ξ^j for $j > i$. We could use a Lévy process, especially one of those that have an expression in terms of a subordinated Brownian motion[3]. Here we decide to use a Constant Elasticity of Variance (CEV) form of local volatility: over the interval $[T_i,\, T_i + \Delta]$ the dynamics of S_t reads:

$$dS = (r_t - q_t)\, Sdt + \sigma_0 \left(\frac{S}{S_{T_i}}\right)^{1-\beta} SdZ_t \tag{3.1}$$

where $\sigma_0\left(\widehat{\sigma}_{VS}\right)$, $\beta\left(\widehat{\sigma}_{VS}\right)$ are functions of $\widehat{\sigma}_{VS} = \sqrt{\xi^i\left(T_i\right)}$ calibrated so that the VS volatility of maturity $T_i + \Delta$ is $\widehat{\sigma}_{VS}$ and the condition on the ATMF skew is fulfilled. r_t and q_t are, respectively, the interest rate and the repo – inclusive of dividend yield. Note that instead of – or in addition to – controlling the skew we could have controlled the convexity of the smile near the money - this would be relevant in the Forex or Fixed Income world. In this article we restrict our attention to the skew.

This completely specifies our model and the pricing algorithm. We can write the corresponding pricing equation as:

$$\frac{dP}{dt} + (r_t - q_t)\, S\frac{dP}{dS} + \frac{\sigma\left(S_{T_{i_0}}, \xi^{i_0},\, S\right)^2}{2} S^2 \frac{d^2P}{dS^2} + \frac{1}{2}\sum_{i,j\,>\,i_0} \rho_{ij}\omega_i\omega_j\xi^i\xi^j \frac{d^2P}{d\xi^i d\xi^j} + \sum_{i\,>\,i_0} \rho_{Si}\omega_i\sigma\left(S_{T_{i_0}}, \xi^{i_0},\, S\right) S\xi^i \frac{d^2P}{dSd\xi^i} = rP$$

where $i_0\left(t\right)$ is such that $t \in [T_{i_0},\, T_{i_0} + \Delta[$, ω_i is the volatility of the ξ^i and ρ_{ij} their correlations.

In the N-factor model, $\omega_i = \omega$ and $\rho_{ij} = \rho\left(Z_i, Z_j\right)$. In the two-factor model, the dynamics of the ξ^i is driven by the processes X and Y. The pricing equation can then be written more economically as:

$$\begin{aligned} &\frac{dP}{dt} + (r_t - q_t)\, S\frac{dP}{dS} - k_1X\frac{dP}{dX} - k_2Y\frac{dP}{dY} + \frac{\sigma\left(\cdots,\, S\right)^2}{2} S^2\frac{d^2P}{dS^2} \\ + &\frac{1}{2}\left(\frac{d^2P}{dX^2} + \frac{d^2P}{dY^2} + 2\rho\frac{d^2P}{dXdY}\right) + \sigma\left(\cdots,\, S\right) S\left(\rho_{SX}\frac{d^2P}{dSdX} + \rho_{SY}\frac{d^2P}{dSdY}\right) = rP \end{aligned}$$

where ρ_{SX} and ρ_{SY} are, respectively, the correlation between Brownian motions U_t and Z_t and the correlation between W_t and Z_t. $\sigma\left(\cdots,\, S\right)$ is a short-hand notation for:

$$\sigma\left(\cdots,\, S\right) \equiv \sigma\left(S_{T_{i_0}}, \xi^{i_0}\left(X_{T_{i_0}},\, Y_{T_{i_0}}\right),\, S\right)$$

4. Pricing

We now turn to using the model for pricing, focussing on the two-factor model. In what follows we choose as time scale $\Delta = 1\,$month. By construction the model is calibrated at time t_0 to the FV curve for all maturities $t_0 + i\Delta$. We specify, in this order:

- values for $k_1,\, k_2, \omega,\, \rho,\, \theta$
- a value for the forward ATMF skew
- values for ρ_{SX} and ρ_{SY}

These steps are discussed in the next three sections.

4.1. Setting a dynamics for implied VS volatilities

Our aim is to price options whose price is a very nonlinear function of volatility; as we roll towards the option's maturity, the maturity of the volatilities we are sensitive to shrinks as well: we thus need the ability to control the term-structure of the volatilities of volatilities, be they ATMF or VS volatilities. In our framework, it is more natural to work with VS volatilities.

In our model the dynamics of VS volatilities is controlled by $k_1,\, k_2, \omega,\, \rho,\, \theta$. As there is presently no active market for options on forward ATM or VS volatility, these parameters cannot be calibrated on market prices. Thus their values have to be chosen so that the level and term-structure of volatility of volatility are reasonably conservative when compared to historically observed volatilities of implied volatilities[4].

[3] For example the Variance Gamma and Normal Inverse Gaussian processes

[4] Dealers trading Napoleons and Reverse Cliquets usually accumulate a negative Gamma position on volatility. In practice, bid and offer term structures of vol–of–vol are used for pricing.

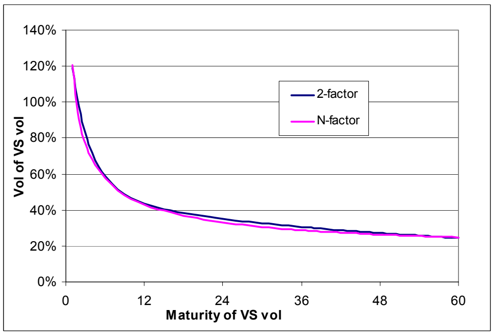

Figure 4.1: Term-structure of the vols of VS vols for a 1 month interval

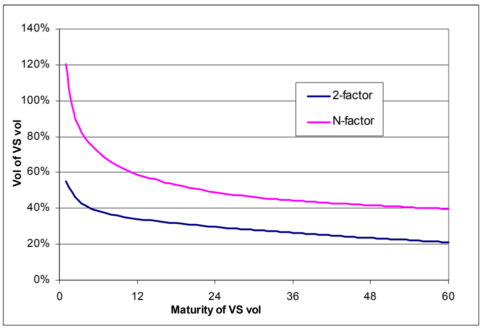

Figure 4.2: Term-structure of the vol of VS vols, for a 1-year interval

Here we choose the following values:

$$\omega = 2.827, \rho = 0, \theta = 30\%, k_1 = 6 \text{ (2 months)}, k_2 = 0.25 \text{ (4 years)} \tag{4.1}$$

so that the volatility of volatility for a one-month horizon is about 120% for the 1-month VS vol, 45% for the one-year vol, and 25% for the 5-year vol. Figure 4.1 displays the term structure of the volatilities of VS volatilities generated by the two-factor model with a flat initial VS term-structure at 20% using these parameter values. We graph

$$\frac{1}{\sqrt{\Delta t}}\mathrm{StDev}\left[\ln\left(\frac{\sqrt{V_{\Delta t}^{\Delta t, \Delta t+\tau}}}{\sqrt{V_0^{\Delta t, \Delta t+\tau}}}\right)\right]$$

for a range of values of τ from 1 month to 5 years. We have picked $\Delta t = 1$ month. The value of ω is chosen so that, over the interval of $\Delta t = 1$ month, the volatility of the one-month VS volatility is 120%.

We also display the term-structure generated by the N-factor model using the following parameters

$$\sigma = 240\%,\ \theta = 40\%,\ \rho_0 = 5\%,\ \beta = 10\%$$

These values are chosen so that, for $\Delta t = 1$ month, the term structure of the two-factor model is matched. Now let us measure volatilities over a time interval of one year, instead of one month (figure 4.2).

They are very different; although both models would yield similar prices for options on VS variances observed one month from now, they would price differently options on VS variances observed in one year. In the two-factor model volatilities of volatilities will tend to decrease as the time scale over which they are measured increases, due to the mean-reverting nature of the driving processes. In the N-factor model, by contrast, they increase: this is due to the fact that forward variances are lognormal – the term-structure would be constant if forward variances were normal.

4.2. Setting the short forward skew

We calibrate the dependence of σ_0 and β to $\widehat{\sigma}_{VS}$, so that the one–month ATMF skew has a constant value - say 5%; we use the 95%–105% skew:

$$\widehat{\sigma}_{95\%} - \widehat{\sigma}_{105\%} \simeq -\frac{1}{10} \left. \frac{d\widehat{\sigma}_K}{d \ln K} \right|_F$$

instead of the local derivative $\frac{d\widehat{\sigma}_K}{d \ln K}$. This defines the functions $\sigma_0\left(\widehat{\sigma}_{VS}\right)$ and $\beta\left(\widehat{\sigma}_{VS}\right)$. This calibration is easily done numerically; we can also use analytical approximations[5] .

If needed, individual calibration of $\sigma_0\left(\widehat{\sigma}_{VS}\right)$ and $\beta\left(\widehat{\sigma}_{VS}\right)$ can be performed for each interval $[T_i, T_i + \Delta]$. Typically the same calibration will be used for all intervals except the first one, for which a specific calibration is perfomed so as to match the short vanilla skew. Here we use the same calibration for all intervals. Figure 4.3 shows functions $\sigma_0\left(\widehat{\sigma}_{VS}\right)$ and $\beta\left(\widehat{\sigma}_{VS}\right)$ for the case of a constant 95%-105% skew equal to 5%.

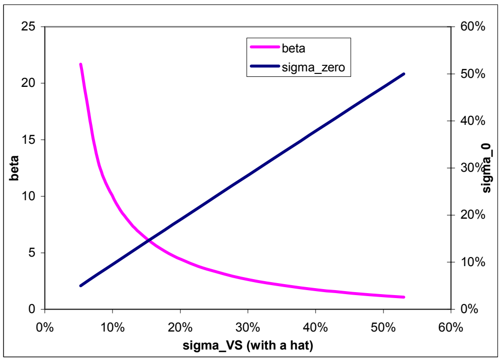

Figure 4.3: β and σ_0 as a function of $\widehat{\sigma}_{VS}$ in the case of a constant 5% skew.

The level of 95%–105% skew can either be selected by the trader or chosen so that market prices of call spread cliquets of period Δ (here one month) are matched.

4.3. Setting correlations between the Spot and short/long factors - the term skew

ρ_{SX} and ρ_{SY} cannot be chosen independently, since X and Y have correlation ρ. We use the following parametrisation:

$$\rho_{SY} = \rho_{SX}\, \rho + \chi \sqrt{1 - \rho_{SX}^2} \sqrt{1 - \rho^2}$$

with $\chi \in [-1, 1]$. ρ_{SX} and ρ_{SY} control both the correlation between Spot and short and long VS volatilities and the term-structure of the skew of vanilla options. They can be chosen, calibrated to the market prices of call spread cliquets of period larger than Δ or calibrated to the market skew for the maturity of the option considered. The dependence of the term skew on ρ_{SX} and ρ_{SY} is made explicit in the following section.

In the N-factor model, we need to specify correlations between the spot process and all forward variances, in a manner which is consistent with correlations specified in equation (2.4), a non-trivial task which we leave outside the scope of this paper.

4.3.1. The term skew

To shed light on how our model generates skew, we derive an approximate expression for the ATMF skew as a function of maturity, for the case of a flat term-structure of VS vols, at order 1 in both ω and the skew $\left. \frac{d\widehat{\sigma}_K}{d \ln K} \right|_F$ at time scale Δ, which we denote Skew_Δ.

Given the skewness $\mathcal{S}_T$ of the distribution of $\ln\left(\frac{S_T}{F_T}\right)$ the ATMF skew is given, at 1st order in $\mathcal{S}_T$ by (Backus, 1997):

$$\text{Skew}_T = \frac{\mathcal{S}_T}{6\sqrt{T}} \tag{4.2}$$

where F_T is the forward for maturity T.

[5] See for example Zhou, (2003)

Consider a maturity $T = N\Delta$ and let us compute the second and third moments of $\ln\left(\frac{S_T}{F_T}\right) = \sum_{i=1}^{N} r_i$ where returns r_i are defined as $r_i = \ln\left(\frac{S_{i\Delta}}{F_\Delta}\right) - \ln\left(\frac{S_{(i-1)\Delta}}{F_{(i-1)\Delta}}\right)$. While returns are not independent, they are uncorrelated. Thus, assuming that Δ is small – so that the drift term in $E[r_i]$ is negligible with respect to the random term:

$$M_3^T = \left\langle \left(\sum_{i=1}^{N} r_i\right)^3 \right\rangle = \sum_i \left\langle r_i^3 \right\rangle + 3\sum_{j>i} \left\langle r_i r_j^2 \right\rangle$$

Let us work at lowest order in Δ: for the purpose of deriving an expression of the third moment at order 1 in ω and S_Δ we can use the following approximations:

$$\begin{aligned} r_j^2 &= \Delta\, \xi^j(T_j) \\ r_i &= \sqrt{\xi^i(T_i)} \int_{T_i}^{T_i+\Delta} dZ_t \end{aligned}$$

Let us denote ξ the constant value of the VS vols at time 0. We get, at order 1 in ω:

$$\begin{aligned} M_3^T &= \sum_i \left\langle r_i^3 \right\rangle + 3\sum_{j>i} \Delta \left\langle \sqrt{\xi^i(T_i)} \int_{T_i}^{T_i+\Delta} dZ_t \xi^j(T_j) \right\rangle \\ &= \sum_i \left\langle r_i^3 \right\rangle + 3\sum_{j>i} \Delta \left\langle \sqrt{\xi^i(T_i)} \int_{T_i}^{T_i+\Delta} dZ_t\, \xi^j(0) \left(\begin{array}{c} 1 + \omega \int_0^{T_j} e^{-k_1(T_j-u)} dU_u \\ + \theta\omega \int_0^{T_j} e^{-k_2(T_j-u)} dW_u \end{array} \right) \right\rangle \\ &= \sum_i \left\langle r_i^3 \right\rangle + 3\sum_{j>i} \Delta\, \omega \xi^j(0) \sqrt{\xi^i(0)} \left\langle \int_{T_i}^{T_i+\Delta} dZ_t \int_0^{T_j} \left(e^{-k_1(T_j-u)} dU_u + \theta e^{-k_2(T_j-u)} dW_u \right) \right\rangle \\ &= \sum_i \left\langle r_i^3 \right\rangle + \rho\omega\xi^{\frac{3}{2}} \Delta^2 N^2 \left[\rho_{SX}\, \zeta(k_1\Delta, N) + \theta\rho_{SY}\, \zeta(k_2\Delta, N)\right] \end{aligned}$$

where $\zeta(x, N)$ is defined by:

$$\zeta(x, N) = \left(\frac{1-e^{-x}}{x}\right) \frac{\sum_{\tau=1}^{N-1} (N-\tau)\, e^{-(\tau-1)x}}{N^2} \tag{4.3}$$

Since we have set the short skew to a value which is independent on the level of variance, expression 4.2 shows that the skewness of r_i is constant. Thus:

$$\sum_i \left\langle r_i^3 \right\rangle = \mathcal{S}_\Delta \sum_i \left\langle \left(\Delta\xi^i(T_i)\right)^{\frac{3}{2}} \right\rangle = N\mathcal{S}_\Delta (\Delta\xi)^{\frac{3}{2}}$$

where $\mathcal{S}_\Delta$ is the skewness at time scale Δ. We then get:

$$M_3^T = N\mathcal{S}_\Delta (\xi\Delta)^{\frac{3}{2}} + \rho\omega\sqrt{\Delta}\, (\xi\Delta)^{\frac{3}{2}} N^2 \left(\rho_{SX}\, \zeta(k_1\Delta, N) + \theta\rho_{SY}\, \zeta(k_2\Delta, N)\right)$$

At order zero in $\mathcal{S}_\Delta$ and ω

$$M_2^T = \left\langle \left(\sum_{i=1}^{n} r_i\right)^2 \right\rangle = N\,\xi\Delta$$

hence the following expression for $\mathcal{S}_T = \frac{M_3^T}{\left(M_2^T\right)^{\frac{3}{2}}}$

$$\mathcal{S}_T = \frac{\mathcal{S}_\Delta}{\sqrt{N}} + \sqrt{N}\omega\sqrt{\Delta} \left[\rho_{SX}\, \zeta(k_1\Delta, N) + \theta\rho_{SY}\, \zeta(k_2\Delta, N)\right]$$

Using eq. (4.2) again we finally get the expression of $\text{Skew}_{N\Delta}$ at order 1 in Skew_Δ and ω:

$$\text{Skew}_{N\Delta} = \frac{\text{Skew}_\Delta}{N} + \frac{\omega}{2} \left[\rho_{SX}\, \zeta(k_1\Delta, N) + \theta\rho_{SY}\, \zeta(k_2\Delta, N)\right] \tag{4.4}$$

This expression is instructive as it makes apparent how much of the skew for maturity T is contributed on the one hand by the instrinsic skewness of the spot process at time scale Δ, and on the other hand by the spot/volatility correlation.

When $\omega = 0$, the skew decays as $\frac{1}{T}$, as expected for a process of independent increments. The fact that volatility is stochastic and correlated with the spot alters this behavior. Inspection of the definition of function ζ in eq. (4.3) shows that for $N\Delta \gg \frac{1}{k1}, \frac{1}{k_2}$, $\zeta(x, N) \propto \frac{1}{N}$, so that $\text{Skew}_{N\Delta} \propto \frac{1}{N}$, again what we would expect.

Equation 4.4 also shows how ρ_{SX} and ρ_{SY} can naturally be used to control the term-structure of the skew.

Figure 4.4 shows how the approximate skew in eq. (4.4) compares to the actual skew. We have chosen the following values: $\Delta = 1\,\text{month}$, the one-month 95%/105% skew is 5%, $\omega = 2.827, \rho = 0, \theta = 30\%, k_1 = 6, k_2 = 0.25$. The spot/vol correlations are: $\rho_{SX} = -70\%$, $\rho_{SY} = -35.7\%$ ($\chi = -50\%$). Even though ω and Skew_Δ are both large, the agreement is very satisfactory.

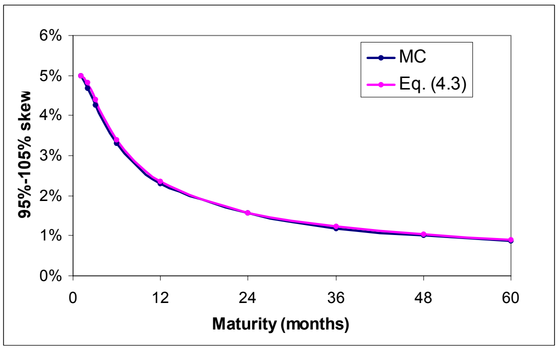

Figure 4.4: The 95%-105% skew as a function of maturity

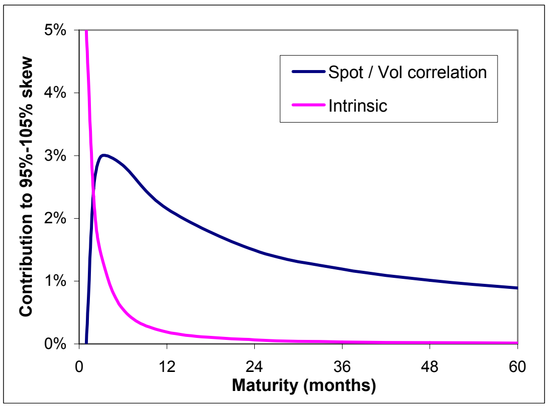

Figure 4.5: The two contributions to the 95%-105% skew in eq. (4.3).

The two contributions to $\text{Skew}_{N\Delta}$ in eq. (4.4) are graphed in Figure 4.5: "intrinsic" denotes the first piece, "spot / vol correlation" denotes the second piece in eq. (4.4). While the contribution of Skew_Δ to $\text{Skew}_{N\Delta}$ is monotonically decreasing, the contribution of the spot/vol correlation is not, as it starts from zero at short time scales. Depending on the relative magnitude of both terms, the term-structure of the skew can be non-monotonic.

While we have derived expression 4.4 for the case of a flat VS term-structure, the general case poses no particular difficulty.

5. Pricing examples

Here we use our model to price a reverse cliquet, a Napoleon, an accumulator, and a call on realised variance, and analyze the relative contribution of forward skew, volatility of volatility and spot/volatility correlation to prices. We use zero interest rates and dividend yield.

For the sake of comparing prices we need to specify how we calibrate model parameters. While it is natural to calibrate to the vanilla smile when pricing options that can be reasonably hedged with a static position in vanilla options, it is more natural to calibrate to call spread cliquets and ATM cliquets when pricing Napoleons and reverse cliquets, which have a large sensitivity to forward volatility and skew.

These products are also very sensitive to volatility of volatility. They are usually designed so that their price at inception is small but increases significantly if implied volatility decreases[6]. As there is as yet no active market for options on variance we use the volatility of volatility parameter values listed in (4.1).

Unless forward skew is turned off, the constant 95%–105% one-month skew is calibrated so that the price of a 3-year 95%–105% one-month call spread cliquet has a constant value, equal to its price when volatility of volatility is turned off and the one-month 95%–105% skew is 5%, which is equal to 191.6%.

In all cases the level of the flat VS vol has been calibrated so that the implied vol of the 3-year one-month ATM cliquet is 20%.

The values for ρ_{SX} and ρ_{SY} are $\rho_{SX} = -70\%$, $\rho_{SY} = -35.7\%$ ($\chi = -50\%$). The corresponding term skew is that of Figure 4.4.

In addition to the BS price, we compute three other prices by switching on either the one-month forward skew ($\widehat{\sigma}_{95\%} - \widehat{\sigma}_{105\%} \neq 0$, $\omega = 0$), or the volatility of volatility ($\widehat{\sigma}_{95\%} - \widehat{\sigma}_{105\%} = 0$, $\omega \neq 0$), or both (full). These prices are listed in table 5.1. We give the definition of each product and comment on pricing results in the following paragraphs.

Model	Reverse Cliquet	Napoleon	Accumulator
Black-Scholes	0.25%	2.10%	1.90%
With Forward Skew	0.56%	2.13%	4.32%
With Vol–of–vol	2.92%	4.71%	1.90%
Full	3.81%	4.45%	5.06%

Table 5.1: A. Model prices

Model	Reverse Cliquet	Napoleon	Accumulator
Full - correlations halved	3.10%	4.01%	5.04%
Full - proportional skew	3.05%	4.30%	4.15%

Table 5.2: B. Model prices

5.1. Reverse cliquet

Here we consider a globally-floored locally-capped cliquet, which pays once at maturity:

$$\max\left(0,\, C + \sum_{i=1}^{N} r_i^-\right)$$

The maturity is three years, returns r_i are observed on a monthly basis ($N = 36$), $r_i^- = \min(r_i, 0)$ and the value of the coupon is $C = 50\%$.

Notice that corrections to the BS price are by no means small, the contribution of volatility of volatility being the largest. The fact that volatility of volatility makes the reverse cliquet more expensive is expected: this option, as well as the Napoleon, is in essence a put on volatility (Bergomi, 2004).

To understand why forward skew increases the price, consider first that, in the four cases listed above, $E\left[\sum_{i=1}^{N} r_i^-\right]$ is constant, by calibration on the ATM cliquet. Next consider the last period of the reverse cliquet: the final payoff is a function of the final return; it is a call spread whose low and high strikes are, respectively, $-C + \sum_{i<N} r_i^-$ – if it is negative – and zero. When forward skew is turned on, the implied vol of the ATM strike is unchanged, by calibration, while the implied vol for lower strikes increases, making the call spread more expensive. The same argument holds for returns prior to the last one.

[6] See Figure 1 in Bergomi (2004).

5.2. Napoleon

The maturity is still three years and the option pays at the end of each year a coupon given by:

$$\max\left(0,\, C + \min_i r_i\right)$$

where r_i are the 12 monthly returns observed each year. Here we use $C = 8\%$.

Again, we notice that volatility of volatility accounts for most of the price. Forward skew seems to have no sizeable impact, however this is not generic; its magnitude and sign depend on the coupon size. While the payoff is still a call spread as a function of the final return, both strikes lie below the money. Also, in contrast to the case of the reverse cliquet, $E\left[\min_i r_i\right]$ is not constant in the four cases considered.

5.3. Accumulator

The maturity is again 3 years with one final payout, given as a function of the 36 monthly returns r_i by:

$$\max(0, \sum_{i=1}^{N} \max\left(\min\left(r_i, \mathrm{cap}\right), \mathrm{floor}\right))$$

where floor $= -1\%$ and cap $= 1\%$ – a standard product.

The largest contribution comes from forward skew. Notice that switching on the volatility of volatility in the case when there is no skew has no material impact on the price while it does when forward skew is switched on. To understand this, observe that, in Black-Scholes, when both strikes are priced with the same volatility, a 99%–101% one-month call spread has negligible vega. However, when the call spread is priced with a downward sloping skew, it acquires positive convexity with respect to volatility shifts.

5.4. Effect of spot/volatility correlation – decoupling of the short forward skew

In standard stochastic volatility models, changing the spot/volatility correlation changes the forward skew and thus the price of cliquets. In our model, because of the specification chosen for the spot dynamics in eq. (3.1), changing the spot/vol correlation does not change the value of cliquets of period 1 month. It only alters the term skew.

Prices quoted above have been computed using $\rho_{SX} = -70\%$, $\rho_{SY} = -35.7\%$. With these values the three-year 95% – 105% skew is 1.25%.

Let us now halve the spot/volatility correlation: $\rho_{SX} = -35\%$, $\rho_{SY} = -18\%$ ($\chi = -19.2\%$). The three-year 95% - 105% skew is now 0.75% – almost halved. The implied vol of the three-year cliquet of one-month ATM calls remains 20% and the price of a 95%-105% one–month call spread cliquet is unchanged, at 191.6% . The corresponding prices appear on first line of table B. The difference with prices on line four of table A measures the impact of the term skew, all else – in particular cliquet prices – being kept constant. The fact that prices decrease when the spot/vol correlation is less negative is in line with the shape of the BS Vega as a function of the spot value[7].

5.5. Making other assumptions on the short skew

Here we want to highlight how a different model for the short skew alters prices, using the three payoff examples studied above. We now calibrate functions $\sigma_0\left(\widehat{\sigma}_{VS}\right)$ and $\beta\left(\widehat{\sigma}_{VS}\right)$ so that, instead of being constant, the 95% – 105% skew for maturity Δ is proportional to the ATMF volatility for maturity Δ.

We have calibrated the proportionality coefficient so that the three-year cliquet of one-month 95% – 105% call spreads has the same value as before. The flat VS vol is still chosen so that the implied volatility of the 3-year cliquet of one-month ATM call is 20%. Prices are listed on the second line of table B.

The accumulator is now sizeably cheaper. One can check that, in Black-Scholes, the value of a symmetrical call spread as a function of ATM volatility, when the skew is kept proportional to the ATM vol, is almost a linear function of volatility – in contrast to the case when volatilities are shifted in parallel fashion, where it is a convex function of volatility – thus explaining why the volatility of volatility has much less impact than in the constant skew case.

[7] See Figure 1 in Bergomi (2004)

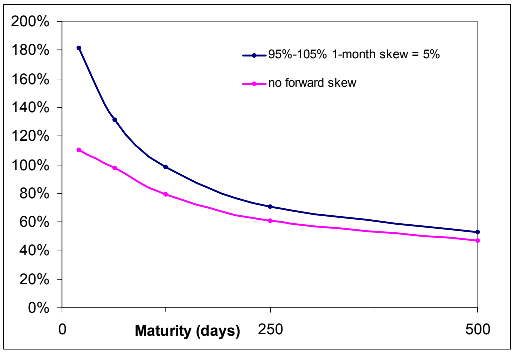

Figure 5.1: Implied volatility of a Call option on realized variance as a function of maturity in the two-factor model

5.6. Option on realized variance

Here we consider a call option which pays at maturity:

$$\frac{1}{2\widehat{\sigma}_K} \max\left(\sigma_h^2 - \widehat{\sigma}_K^2, 0\right)$$

where volatility $\widehat{\sigma}_K$ is the strike, and σ_h^2 is the annualized variance measured using daily log-returns. We assume there are 250 daily observations in a year, equally spaced. The variance of the distribution of σ_h^2 has two sources:

- The dynamics of VS variances
- The fact that observations are discrete: in the case when VS variances are static it is the only contribution and it is determined by the distribution of spot returns, in particular its kurtosis, which depends on assumptions made for the short-maturity smile – in our context, the value of β. In the general case, it affects short-maturity options most noticeably.

Prices are expressed as implied volatilities computed with the Black-Scholes formula with zero rate and repo. The underlying is the VS variance for the maturity of the option, whose intial value is given by the VS term-structure observed at the pricing date.

In our model, daily returns are generated by the volatility function form in eq. (3.1). Their conditional kurtosis is a function of β, a parameter we use to control the short-term skew. The prices of options on variance will thus depend on assumptions we make for the skew at time scale Δ. Figure 5.1 shows implied vols of call options on Variance, using a flat term-structure of VS vols at 20%, the same correlations as in the examples above ($\rho_{SX} = -70\%$, $\rho_{SY} = -35.7\%$) , for the two cases: $\widehat{\sigma}_{95\%} - \widehat{\sigma}_{105\%} = 5\%$ and $\widehat{\sigma}_{95\%} - \widehat{\sigma}_{105\%} = 0\%$.

Figure 5.1 illustrates how assumptions for the forward skew significantly affect the distribution of returns and thus the price of options on variance, mostly for short options. The shortest maturity in the graph corresponds to options of maturity 1 month (20 days): since we have taken $\Delta = 1$ month, the distribution of σ_h^2 does not depend on the dynamics of variances ξ^i – it only depends on β.

Note that, in this model, VS volatilities for maturities shorter than Δ are not frozen: instead of being driven by eq. 2.3, their dynamics is set by the value of β.

5.6.1. Using the N-factor model

It is instructive to compare prices of options on realised variance generated by the two-factor and N-factor models. As figures 4.1 and 4.2 illustrate, even though the short-term dynamics of VS vols in both models are similar, they become different for longer horizons.

Here we price the same option on variance considered above using the N-factor model of forward variances. Parameter values for the dynamics of forward variances are the same as those used in figures 4.1 and 4.2. We have taken no forward skew: ($\widehat{\sigma}_{95\%} - \widehat{\sigma}_{105\%} = 0$). Also, to make prices comparable with those obtained in the two-factor model, we have taken zero correlation between spot and forward variances. Implied volatilities for both models are shown in figure 5.2.

Because $\Delta = 1$ month, for thave he shortest maturity considered – 20 days – the implied volatilities for both models coincide. For longer maturities the fact that implied volatilities are higher in the

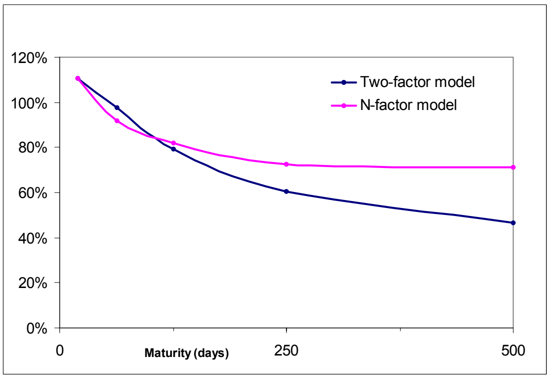

Figure 5.2: Implied volatility of a Call option on realized variance as a function of maturity in the two-factor and N-factor models

N-factor model is in agreement with the graph in 4.2 which shows that, for longer horizons, the volatilities of forward variances in the N-factor model are larger than in the 2-factor model.

Finally, in addition to the effects discussed above, prices of calls on variance have to be adjusted to take into account bid/offer spreads on the VS hedge. These can be approximately included by shifting the level of volatility of volatility (Leland, 1985).

6. Conclusion

We have proposed a new model that, in contrast to popular stochastic volatility and Jump/Lévy models, gives us the flexibility to independently control:

- the term-structure of the volatility of volatility,
- the short-term skew,
- the correlation of spot and volatilities.

This model consistently prices general exotic options and options on variance or volatility. We achieve this by choosing a time scale Δ, specifiying how the forward skew for maturity Δ depends on the level of volatility, and modelling the dynamics of VS variances. The model can be simultaneously calibrated on the short vanilla skew, the long vanilla skew, an ATM cliquet and a call spread cliquet of period Δ, while letting us specify freely the dynamics of forward variances as well as how the short forward skew depends on volatility.

By directly controlling the short forward skew we are able to adjust the amount of skewness of the distribution of $\ln(S_T)$ generated on one hand by the intrinsic skewness of the process at short time scales, and the spot/volatility correlation on the other hand. Handling this issue appropriately is, in our view, an essential task in the design of models which accurately capture the three effects mentioned above.

Even though the choice of the time scale Δ is natural for many payouts – for example Napoleons and reverse cliquets – it is more arbitrary for other options, for example options on variance. It would be useful to have scaling relationships relating parameter sets for different values of Δ so that some model features remain unchanged – for example the skew of vanillas or the implied forward skew of cliquets. This is left for future work.

As of today there are no market prices for caps/floors/swaptions on forward variances. Choosing a value for the parameters governing the dynamics of forward variances is thus a trading decision. It is the hope of the author that liquidity of options on volatility and variance increases so that we will soon be able to trade the smile of the volatility of volatility!

About the author

Lorenzo Bergomi is head of Quantitative Research in the Equity Derivatives department at Société Générale. He wishes to thank members of his team for useful discussions.Email: lorenzo.bergomi@sgcib.com

References

Backus D., Foresi S., Li K. and Wu L., 1997
Accounting for biases in Black-Scholes
Unpublished

Bergomi L., 2004
Smile dynamics
Risk September, pages 117-123

Dupire B., 1996
A unified theory of volatility
Unpublished

Leland H., 1985
Option replication with transaction costs
The Journal of Finance, 40 (5), pages 1283-1301

Zhou F., 2003
Black smirks
Risk January, pages 87-91